

OPENCORVUS: EVOLVING EXPERT ORGANIZATIONS FOR PROFESSIONAL WORK

Author identities pending

Non-experimental working manuscript

ABSTRACT

Professional agent systems must carry specialized procedures across assignments and improve those procedures without losing the ability to identify what was executed and evaluated. We introduce Evolving Expert Organizations (EEOs), a framework that makes a reusable expert organization the unit of specialization and revision. Implemented in OpenCorvus, an organization packages role instructions, collaboration graphs, capability allocations, domain resources, and exact version identity. Domain specialization binds this structure to professional procedures and output contracts. An outer improvement process attributes failures from task evidence, constructs candidate revisions, and compares them with their parent under matched execution conditions. We define the admissible organizational update space and instantiate it with experience-driven procedural text revision while preserving executable resources and inherited capability references. Task-local execution remains bound to its original revision, and retained updates apply to subsequent work. Our analysis connects organizational improvement to search-cost amortization and independent evaluation after selection. A controlled experimental design separates expert collaboration, domain specialization, and retained updating, with total-budget accounting and capability-retention measurements. The framework provides a concrete basis for investigating how professional agent organizations accumulate useful procedures across tasks.

1 INTRODUCTION

Professional work is organized around persistent responsibilities: a researcher collects evidence, an analyst transforms it, a reviewer checks the resulting claims, and a writer produces a deliverable. The organization carries procedures, tools, and handoff conventions from one assignment to the next. For a large language model (LLM) system, reproducing this arrangement requires more than assigning expert names. An analyst needs the correct computation tools and input resources; a reviewer needs the actual revision under review; and improvements discovered during one assignment must survive beyond its conversation.

We introduce *Evolving Expert Organizations* (EEOs), a framework for representing, specializing, executing, and revising reusable agent organizations. Its implementation in OpenCorvus treats an expert organization as a versioned package containing role instructions, capability allocations, task-level collaboration graphs, and domain resources. A task binds to an exact package revision. An outer improvement process constructs a new revision from recorded experience and compares it with its parent on matched tasks. The resulting object of reuse is the organization, including its procedures and executable resource closure.

Consider a research report whose calculations are correct but whose prose uses an earlier result. Repairing that report addresses one assignment. Refining the writer’s procedure so that each numerical claim is explicitly bound to its canonical result and revision-specific review targets a recurring organizational failure. The latter is an organizational update only when it is retained in a package and applied to subsequent tasks. Whether it improves those tasks is a separate, measurable question. This distinction connects the representation, update mechanism, and evaluation throughout the paper.

The paper makes three contributions. First, we define a portable organization representation that separates reusable expertise from task-local state and gives domain specialization an explicit resource and output contract. Second, we define a constrained revision protocol that connects the mutation

surface, matched parent–candidate evaluation, and retained version-specific evidence. Third, we derive the budget and selection conditions needed to assess improvement on future work and specify a controlled evaluation that separates the effects of multiple experts, domain specialization, and organizational updating. Our technical analysis concerns these mechanisms; the empirical protocol evaluates their effect on quality, cost, and retained capabilities.

2 RELATED WORK

Expert collaboration. Role specialization and structured communication are established design choices. AutoGen supplies configurable conversational agents, and MetaGPT uses professional roles and operating procedures to coordinate software work (Wu et al., 2023; Hong et al., 2023). EEO makes the portable package the unit of specialization and revision: roles are coupled to their capability references, supporting resources, and handoff declarations. Neither the use of expert roles nor a dependency graph alone constitutes the contribution.

Automatic agent and workflow design. EvoAgent mutates agent settings, selects diverse offspring with an LLM quality check, and integrates their task outputs (Yuan et al., 2025). AFlow represents workflows in code and uses execution feedback in a variant of Monte Carlo tree search (Zhang et al., 2025). ADAS uses a meta agent to program new agents from an archive of prior designs and evaluations, including transfer across tasks, domains, and models (Hu et al., 2025). These results motivate treating organizational revision as search, but they preclude claiming that automatic design or cross-task reuse originates here. EEO studies a constrained package search space that carries domain resources and capability allocations together with instructions. Its current generator edits the procedural layer; broader structural edits are represented and checked by the package integrity interface.

Reflective prompt optimization. GEPA (Genetic-Pareto) already uses execution feedback to revise module prompts, retain candidates, and optimize held-out performance with fixed model weights (Agrawal et al., 2025). EEO’s text configuration therefore does not introduce reflection-based optimization. Its contribution concerns professional organization packaging, explicit resource boundaries, and attributable revision retention. Section 6 includes a matched optimizer comparison on this substrate.

Learning from experience. SiriuS retains successful reasoning trajectories, augments unsuccessful ones, and performs supervised fine-tuning of individual agent models (Zhao et al., 2025). EEO holds the foundation model fixed and retains improvements in organization packages. The two update different objects and can be combined, but an experience archive by itself does not establish learning: a retained package change and its measured transfer are the relevant evidence here.

Execution substrate. Agent environments and interfaces, exemplified by SWE-agent and OpenHands, determine how agents interact with tools and artifacts (Yang et al., 2024; Wang et al., 2024). The OpenHands Software Agent SDK also provides persistent event-oriented execution (Wang et al., 2025). EEO uses exact package binding and revision-specific artifacts to keep organizational comparisons attributable. These supporting mechanisms do not by themselves imply better reasoning or semantic correctness.

3 EXECUTABLE EXPERT ORGANIZATIONS

3.1 REPRESENTATION AND EXECUTION

An organization revision is

$$\mathcal{O}_v = (R_v, G_v, P_v, K_v, A_v, F_v; h_v). \quad (1)$$

R_v is the set of expert roles. G_v is a collection of directed acyclic graphs whose nodes name roles and whose edges declare task-level dependencies. P_v contains the orchestrator and expert instructions. K_v contains model-readable skills, procedures, references, and examples. A_v maps the orchestrator and roles to declared capability references, including tools and skills. F_v is the complete package file

set, containing these resources and any executable or opaque assets. The content digest h_v identifies the exact revision; a human-readable version label is additional metadata.

A role may occur in different workflow declarations, so roles and workflow nodes are distinct objects. The graphs expose intended dependencies to the orchestrator. Dispatch remains an LLM decision using predecessor evidence, active workers, and execution capacity; graph validity is not a proof that every dependency is semantically satisfied at runtime.

For task x , execution produces a trajectory τ and deliverable y :

$$(\tau, y) \sim \text{Run}_\phi(\mathcal{O}_v, x, \eta, b). \quad (2)$$

Here ϕ fixes the foundation model and inference configuration, η specifies the workspace, tools, permissions, and external environment, and b is the task budget. The task stores an immutable package binding. Worker descriptors carry the corresponding projected capabilities and revision identity. Conversation state, intermediate artifacts, and scheduler ownership belong to the execution, rather than to the reusable organization.

Figure 1 shows the representation through the Research Studio package. Its five expert roles cover planning, research, analysis, fact checking, and writing. The analyst produces canonical computation evidence before fact checking; the writer renders the accepted evidence rather than recomputing numerical claims. These are declared responsibilities and workflow relationships in the inspected package (OpenCorvus contributors, 2026).

The same package declares three workflows with different entry conditions. Full research uses all five roles when the assignment requires new source collection. Evidence synthesis starts with the analyst when a sufficient source bundle is already supplied. Direct writing uses the writer alone for bounded material that requires faithful presentation without new research claims. These alternatives reuse the same role definitions and resources; they do not require a separately maintained organization for each task shape. Choosing among them remains part of the orchestrator’s task interpretation. A task with missing computational evidence cannot become a direct-writing task merely by asking the writer to supply the missing calculations.

3.2 DOMAIN SPECIALIZATION

We express a domain specification as $\mathcal{D} = (\mathcal{X}, K^\mathcal{D}, A^\mathcal{D}, C^\mathcal{D}, q^\mathcal{D})$: a task family, domain procedures, required capabilities, output contracts, and an evaluation rubric. A specialization operation

$$\mathcal{O}_0^\mathcal{D} = \text{Specialize}(\mathcal{O}_{\text{base}}, \mathcal{D}) \quad (3)$$

constructs a package with explicit role responsibilities and resource bindings. $C^\mathcal{D}$ and $q^\mathcal{D}$ are analytical names for the domain’s artifact schemas, checking procedures, and scorers; they are not additional mandatory manifest fields.


For evidence-based analysis, specialization assigns collection and citation procedures to the researcher, executable calculations and canonical result schemas to the analyst, verification criteria to the fact checker, and a report template to the writer. The evaluation checks numerical agreement with the canonical results, claim support, and delivery requirements. Replacing the names with professional titles while keeping tools, procedures, and acceptance criteria unchanged is an insufficient specialization.

Initial specialization can introduce domain resources and authorized tools. Subsequent evolution operates within the parent package’s inherited capability envelope and frozen executable resources. This separates establishing what an organization can access from learning how it uses those resources. It also makes the comparison with a fixed specialized organization meaningful: both start with the same professional substrate.

3.3 A CONCRETE PROFESSIONAL CONTRACT

The analytical-report specialization gives organizational knowledge three connected forms. A *procedure* explains how to perform a recurring operation, such as reconciling a source population before aggregating it. A *resource binding* makes the corresponding input, tool, schema, or reference available to the responsible role. An *artifact contract* describes the evidence a downstream role

(a) Reusable organization



(b) One package, task-local executions




Figure 1: A reusable organization and its task-specific workflow choices. The five expert roles and three alternative workflow declarations follow Research Studio. The image abbreviates the package digest h_v as h . Solid role-to-role arrows are declared dependencies; dashed orchestrator links indicate LLM dispatch. The resource bands belong to the package, while outputs and execution state belong to individual tasks.

needs before using the output. These forms have different failure modes: good instructions cannot compensate for an unavailable source, and an available calculation tool does not specify which population or units are appropriate.

For a report task, the planner fixes the question, source hierarchy, comparison dimensions, and stopping conditions. The researcher assembles durable sources and identifies disagreement. The analyst records the calculation method, source identities, canonical tables, and claim-to-result correspondence. The fact checker examines those post-computation claims and records corrections and unresolved limits. The writer constructs a structured report model from the accepted material and derives its presentation from that model. This allocation follows the Research Studio instructions (OpenCorvus contributors, 2026). It places the numerical method upstream of verification and presentation, so a persuasive final report is not the only object available for inspection.

To make the handoff precise, consider a quantitative claim κ . Its analytical support record can be written as $(\kappa, r, a, \ell_{\text{rev}})$, where r identifies a canonical result revision, a locates the relevant table cell or range, and ℓ_{rev} identifies the review addressing that claim and result. This notation describes the information relationship, not a new mandatory runtime schema. A complete record still needs semantic checking: the table may use the wrong population, the cited range may not entail the prose, or the review may leave the claim unresolved. The writer’s contract therefore concerns both locating the resources and preserving the accepted meaning, units, qualifications, and uncertainty.

The reusable package contains the rules for constructing and checking such relationships. Task-specific sources, computed values, and accepted conclusions remain in the task’s evidence. Thus, specializing an organization for analytical reporting does not require installing the answer to each report in its prompts. It establishes a procedure that can be applied to different questions, source bundles, and result tables. Whether this procedure transfers well is measured on new tasks with separately constructed evidence.

4 EXPERIENCE-DRIVEN ORGANIZATIONAL EVOLUTION

The improvement process runs across tasks. A new conversation, additional worker, or repaired output within Equation 2 does not change $\mathcal{O}_v$. Organizational evolution produces and retains a distinct package revision for subsequent tasks.

4.1 EXPERIENCE AND ATTRIBUTION

We represent the evidence associated with an evaluated execution as

$$e = (x_{\text{id}}, h_v, \gamma, \ell_\tau, \ell_y, s, c, u, \iota), \quad (4)$$

where γ identifies the frozen evaluation context; ℓ_τ and ℓ_y locate run evidence and deliverable resources; s is a scorer vector; c records resource consumption; u records human intervention; and ι locates independent integrity review. This tuple joins the implementation’s separate run, evaluation, resource, and review records; intervention time is an additional study measurement. Task and package identities make the evidence attributable to a particular revision.

The Evolution Lab divides improvement work among observation, failure analysis, experiment planning, candidate authoring, evaluation, integrity review, and recommendation roles. The failure analyst identifies a recurring procedural mismatch and links it to concrete runs. For the running example, the hypothesis is that obligations distributed across the writer instructions and quality skill need a more explicit local sequence where individual claims enter the report. A candidate elaborates the claim-to-result handoff procedure and its reference example. This attribution is a testable hypothesis, not a causal conclusion inferred solely from an LLM explanation.

4.2 CANDIDATE SPACE AND UPDATE POLICY

A candidate $\mathcal{O}'$ must preserve the package’s logical identity, declare a new version, and contain a substantive change. Let $\text{Cap}(\mathcal{O})$ denote the union of locally expanded capability references declared by its orchestrator and agents; let $\text{Base}(\mathcal{O})$ denote their base-role set. Same-package capability sets are expanded to member references; external references retain their encoded identity. The integrity relation requires

$$\text{Cap}(\mathcal{O}') \subseteq \text{Cap}(\mathcal{O}_v), \quad \text{Base}(\mathcal{O}') \subseteq \text{Base}(\mathcal{O}_v), \quad (5)$$

$$F'|_{\overline{M}} = F_v|_{\overline{M}}, \quad M = M(\mathcal{O}_v) \cup M(\mathcal{O}') \cup \{m\}. \quad (6)$$

$M(\mathcal{O})$ contains declared prompts, the package README and selector instructions, and eligible UTF-8 skill instructions, references, and examples. Here m denotes the manifest path, handled as a mutable declaration. Equation 6 requires both file-set equality and byte equality outside the two text closures and the manifest. Candidate validation additionally resolves declared prompt paths, workflow agents, and dependencies, and rejects cycles.

The representation admits four kinds of change: procedural text; redistribution of inherited capabilities; workflow dependencies; and the role roster. Figure 2 develops a procedural candidate within a concrete professional specialization. The union constraint in Equation 5 is package-wide. A newly introduced role can receive a capability previously held elsewhere in the parent, so it is not a per-role non-escalation guarantee.

The released candidate author is instructed to edit procedural text:

$$\mathcal{O}' \sim \mu_\phi(\cdot \mid \mathcal{O}_v, \mathcal{E}_{\text{dev}}, H), \quad (7)$$

where H is the attributed failure hypothesis. The intended text space $\mathcal{S}_{\text{text}}(\mathcal{O}_v)$ fixes the role roster, topology, and capability allocations while editing eligible instructions and references. A prompt does not guarantee that every sampled proposal belongs to this space. The research arm therefore checks the realized diff against $\mathcal{S}_{\text{text}}$ before admitting a candidate to its trials; rejected attempts still consume its budget. This arm-eligibility check is part of the experimental protocol, distinct from the broader production validator. Structural mutation is an extension supported by the representation and integrity interface, rather than an automatically explored operator in this configuration.

4.3 WORKED EXAMPLE: REFINING A CLAIM HANDOFF

Consider a development case in which an analyst publishes a corrected result after an earlier calculation, and a report incorporates a value from the earlier revision. The arithmetic in the corrected resource is sound, but the final numerical claim has the wrong evidential basis. This is an illustrative execution pattern for explaining candidate construction, not a reported experimental case. The current writer instructions already require exact analysis and review resources. The shared quality skill additionally requires claim identifiers, result-table coordinates, verification evidence, and a report model populated from accepted claims. The candidate therefore changes where and in what order these existing obligations are presented to the writer; it does not introduce the evidence contract itself.

The failure analyst first distinguishes competing explanations. The writer may have read an old resource; it may have read the correct resource but transcribed the wrong cell; or the report renderer may have consumed a stale report model. These explanations require different revisions. A run whose locators identify the accepted analysis but whose report-model claim points to a superseded table supports investigating the handoff procedure. A run whose model is correct but whose rendered document differs instead directs attention to the rendering path. The evidence narrows the candidate hypothesis; a fluent retrospective explanation alone does not identify the cause.

For the handoff hypothesis, the text candidate reorganizes those distributed obligations into a checklist placed at the writer’s claim-construction step: open the accepted result; resolve the associated review; bind the claim to its table coordinate; and write the value together with its qualifications and unresolved limits. A companion example contrasts a corrected result with a superseded one using generic placeholders. This changes the presentation and point of use of the existing requirements, while preserving their underlying acceptance contract. The update retains the role roster, collaboration declarations, capabilities, and executable files. Its reusable content is the sequence of checks and the reference pattern, rather than a copied task answer or a task-specific source locator. Figure 2 shows this relationship between specialization, evidence, and procedural revision.

The transfer hypothesis concerns the effect of that reorganization: tasks that require carrying revised quantitative evidence into a report should exhibit fewer result–report mismatches. The same procedure can also add retrieval and checking cost, cause excessive abstention when review evidence is partial, or distract from qualitative synthesis. Matched trials therefore examine contract satisfaction and resource use, and retention tasks include assignments for which the extra numerical procedure is irrelevant. If improvement occurs only on the development example, the evidence supports fitting that example rather than transferable organizational learning. If output correctness remains unchanged while checking cost rises, the revision can be valid and precisely attributable without being useful.

4.4 MATCHED COMPARISON AND RETENTION

A campaign fixes task cases, repetitions, scorer digests, model configuration, environment, permissions, arm order, and resource ceilings. Baseline and candidate executions occupy matched case–repetition slots and use their exact package revisions. The publisher derives candidate identities and file differences from validated resources. Evaluation results are projected from metric execution receipts, and independent integrity review binds to the exact evaluated revision.

For scorer j , let s_{ij}^a be the value for slot i and arm $a \in \{0, 1\}$, with 0 denoting the parent. The comparison computes

$$d_{ij} = s_{ij}^1 - s_{ij}^0, \quad \bar{d}_j = \frac{1}{n} \sum_{i=1}^n d_{ij}, \quad \hat{\Delta} = \sum_j w_j \frac{\sigma_j \bar{d}_j}{|t_j - f_j|}. \quad (8)$$

Weights w_j are the campaign’s nonnegative weights normalized to sum to one, $\sigma_j \in \{-1, 1\}$ sets the preferred direction, and distinct target and floor values t_j, f_j define the scale. Missing required runs, scorers, or reviews make a comparison unavailable rather than contributing a zero score. The implementation reports paired summaries and an operational retain/select/inconclusive recommendation. Its multi-scorer interval approximation is discussed in Appendix A; confirmatory inference uses the separate task-level analysis in Section 5.

Retention creates a traceable successor relation $\mathcal{O}_v \rightarrow \mathcal{O}_{v+1}$ and records the comparison context. Activation applies the selected revision to new tasks, preserving existing task bindings. A user-requested revision can also be retained through explicit feedback, but that path carries no claim of




Figure 2: Domain resources and professional responsibilities support a shared analytical contract. The worked candidate reorganizes obligations already distributed across the writer prompt and shared quality skill into an explicit sequence at claim construction (Section 4.3). The transfer hypothesis concerns this procedural reorganization; it does not establish a new evidence contract. The fixed-update footer describes the EEO-text research arm.

measured improvement. Selection, installation, and demonstrated learning therefore have distinct evidence.

Table 1 summarizes one complete improvement cycle. Candidate generation and matched comparison can repeat within the frozen development budget; the independent test partition is opened only after the final selection.

Step	One organizational improvement cycle
1	Freeze the parent package, development inputs, scorer versions, and total search budget.
2	Inspect development run evidence; attribute a recurring failure and state a candidate hypothesis.
3	Generate a text-policy candidate; check the realized diff for text-arm eligibility, then validate the package, inherited capabilities, and frozen files. Charge rejected attempts to search.
4	Execute matched parent and candidate trials; collect metrics, failures, resource use, and independent integrity review.
5	Record the comparison. Retain the parent or select an eligible candidate; keep each evaluated revision and its evidence addressable.
6	Freeze the final selected revision before evaluation on unseen tasks and capability-retention tasks. Start later development with a fresh evaluation partition.

Table 1: EEO improvement procedure. Steps 1–5 use Evolution Lab’s package and evaluation interfaces. Step 6 specifies the research evaluation boundary; reuse of a test partition for subsequent candidate selection invalidates that boundary.

5 ANALYSIS: WHAT CONSTITUTES IMPROVEMENT?

5.1 RESOURCE CLOSURE AND REVISION ISOLATION

The inherited-reference constraint composes across retained revisions. If every transition satisfies Equation 5, then, by transitivity of set inclusion, $\text{Cap}(\mathcal{O}_v) \subseteq \text{Cap}(\mathcal{O}_0)$ and likewise for base roles. Together with the frozen-file comparison, this bounds the resources an accepted candidate can introduce through this update path. It does not imply semantic safety: a changed instruction can misuse an existing tool, and a remote tool or externally resolved capability set can change without changing its reference.

Exact task binding also separates an update’s intended unit of effect. Executions already bound to h_v continue to identify $\mathcal{O}_v$; future tasks bound to h_{v+1} identify the successor. This removes one source of attribution ambiguity, but it does not make the external environment stationary. Comparison contexts must still account for tool versions, model configuration, workspace state, and nondeterminism.

5.2 TOTAL VALUE AND AMORTIZATION

Let $q(\mathcal{O}, x) \in [0, 1]$ measure task quality and let $c(\mathcal{O}, x)$ and $u(\mathcal{O}, x)$ measure execution cost and intervention time. For fixed nonnegative conversion weights λ, ρ , define per-task utility

$$U(\mathcal{O}, x) = q(\mathcal{O}, x) - \lambda c(\mathcal{O}, x) - \rho u(\mathcal{O}, x). \quad (9)$$

Quality, cost, and time are also reported separately; the weights express an application-specific tradeoff rather than an intrinsic quality measure.

Suppose a revision has expected utility gain g over its parent on the future task distribution, and the entire search and validation process costs C_{search} in the same utility units. Over T tasks, its expected net advantage is

$$V_T = Tg - C_{\text{search}}, \quad V_T > 0 \iff T > C_{\text{search}}/g \quad (g > 0). \quad (10)$$

This accounting identity explains why an improved task score is insufficient evidence of useful evolution. Search includes unsuccessful candidates, experiment planning, both trial arms, evaluator calls, and integrity review. When $g \leq 0$, additional reuse cannot repay a nonnegative search cost under this stationary model. When tasks change over time, the relevant expression is $\sum_{t=1}^T g_t - C_{\text{search}}$.

5.3 SELECTION AND UNSEEN-TASK EVIDENCE

A development winner is selected partly because of favorable observations. We therefore evaluate the frozen winner and parent on tasks withheld from candidate construction and selection. Let $d_i = q(\mathcal{O}', x_i) - q(\mathcal{O}_v, x_i) \in [-1, 1]$ be the paired quality difference on independent held-out task units. Each d_i can average multiple repeated executions within the task; those repeats do not increase the number of independent tasks.

Conditional on the fixed organizations, Hoeffding’s inequality gives (Hoeffding, 1963)

$$\Pr(|\bar{d} - \mathbb{E}\bar{d}| \geq \epsilon) \leq 2 \exp(-n\epsilon^2/2), \quad \epsilon(n, \alpha) = \sqrt{\frac{2 \log(2/\alpha)}{n}}. \quad (11)$$

This conservative bound remains valid for nonidentical independent task distributions, with $\mathbb{E}\bar{d}$ interpreted as their average effect. A claim about a target population additionally requires representative sampling. Repository or source-document dependence requires defining the independent unit at that cluster level.

If K candidates are all fixed before inspecting the same held-out set, a union bound replaces $\log(2/\alpha)$ with $\log(2K/\alpha)$ for simultaneous intervals. This finite-family bound does not justify adaptively generating new candidates from held-out results. Such feedback converts the set into development data and requires a fresh test set. At $K = 1$ and $\alpha = 0.05$, a worst-case half-width of 0.10 requires $n \geq 738$ independent units; a small pilot cannot establish a precise population-level improvement. The inequality is a standard statistical tool, not a new learning guarantee. Its estimand is the effect of the frozen revision pair. It does not average over the random search process that produced the candidate.

6 EXPERIMENTAL DESIGN

The evaluation targets four questions: whether specialization improves professional task quality; whether retained updates improve unseen work; whether any gains justify total resources and human effort; and whether improvement preserves previously useful capabilities. This manuscript contains the method and experimental design; performance measurements are not yet reported.

Task families and partitions. The primary family is evidence-based analytical reporting: each task provides a fixed source bundle, a computational question, and a deliverable contract. Acceptance requires reproducible calculations, support for load-bearing claims, and agreement between the rendered report and canonical results. A second family uses software maintenance tasks with executable functional checks and exact repository snapshots. Enrollment requires at least three dependent stages, an intermediate artifact consumed and checked in a later stage, and an independently checkable final deliverable. These requirements apply to the task contract, not the number of agents. We record required stages, tool interactions, and input resource size. The reuse horizon T counts subsequent assignments to the retained organization; it does not assert a particular single-task duration or context length.

Tasks are partitioned by source provenance or repository, before development, into update tasks, unseen future tasks from the same family, and a retained-capability test set. Reworded questions over the same evidence and related issues from one repository remain in the same partition. A separate, untouched family tests cross-domain transfer only when its tools and contracts are already compatible with the starting package; results from this setting are reported separately from within-domain generalization. Figure 3 specifies the information boundary.

Controlled comparisons. All arms use the same foundation model, inference configuration, domain input access, and total resource ceiling. The comparisons in Table 2 isolate successive organizational changes. The strong single agent receives the same available domain resources and can plan, call tools, check results, and revise within its budget. It is not restricted to a one-shot prompt.

Arm	Organization	Domain procedures	Retained updates
Single agent	One iterative tool-using agent	Same resource access	None
Fixed team	Generic expert roles	Generic role procedures	None
Specialized team	Fixed domain organization	Explicit procedures and contracts	None
EEO-text	Same initial specialized team	Same initial package	Experience-driven text revision

Table 2: Primary comparisons. EEO-text is the implemented generator configuration. Structural search is evaluated as a distinct extension only with a specified generator and completed executions; it is not silently included in this arm.

Search costs are charged before future execution. Each arm receives the same episode-wide ceiling B over the same number of assigned future tasks. Before inspecting any test outcome, each arm allocates its post-search remainder to fixed per-task caps whose sum does not exceed that remainder. An arm with no search can allocate larger inference caps. Caps are not recycled across test tasks, and each execution starts from its prescribed isolated workspace. This prevents a costly early task from reducing a later task’s allowance. Unfinished assigned tasks remain in the quality denominator. A supplementary comparison fixes per-task inference cost to expose the mechanism’s effect without amortization; it is labeled separately and reports search overhead.

Optimizer control. The protocol also compares EEO-text with a GEPA-based adaptation on the same initial package, eligible text files, feedback access, model, and total budget (Agrawal et al., 2025). Each eligible file is an optimization module; valid updates are reconstructed as package revisions. Both methods use identical integrity and test procedures. This adapted baseline requires a frozen implementation before execution. It distinguishes revision-policy value from the benefit of updating at all.




Figure 3: Organizational development and independent evaluation. Development trials supply evidence for candidate construction and retention. The final selected organization is frozen before independent unseen-task and prior-capability evaluation. The pairing matrix denotes trial slots, not measured results. Total resource accounting includes failed search attempts. A new search episode uses fresh partitions; test outcomes do not return to the same search.

Outcomes and scoring. The primary outcome is binary contract satisfaction at the task level. Reporting tasks require computational checks and blinded review against a frozen rubric; software tasks require the prescribed functional checker. An LLM judge is supplementary and uses a separately fixed prompt and model. We report inter-reviewer agreement, adjudication rules, and task-level disagreement counts. Secondary outcomes are component quality, all model and tool consumption, elapsed time, intervention count and minutes, completion rate, and retention-set quality. Missing measurements, infrastructure failures, agent failures, and budget exhaustion receive explicit categories. Agent failure and budget exhaustion count as unsuccessful assigned tasks; infrastructure exclusions follow a rule fixed before arm outcomes are inspected.

Mechanism studies. The attribution ablation supplies both authors with identical raw run evidence, scores, and evaluator feedback. Both arms compute the analyst’s structured hypothesis H and charge its cost; only the treatment author receives H . Candidate-author calls and trial budgets are matched. This isolates the additional attribution message from access to the original evidence. A second ablation removes reusable domain references while preserving role instructions and tools. These tests distinguish the value of experience and resource packaging from generic extra inference. Role, dependency, and capability-allocation changes require separate structural-search arms. The EEO-text arm checks each realized candidate diff for unchanged roles, dependencies, capabilities, and non-version manifest fields; out-of-arm proposals count as rejected search attempts. Validity checks remain active in every arm.

Search-level replication. A single search followed by many test tasks estimates the effect of one retained revision pair. Method-level comparisons, including GEPA, use independent search episodes with fresh development partitions and search randomization. Each episode freezes its candidates and evaluates all arms on a common, previously untouched task set; source clusters are disjoint across episodes. It has its own total budget and fixed test caps. The episode’s paired mean quality difference is one observation for the method comparison. Equation 11 then uses the number of independent

episodes, not the pooled test-task count. Within-episode task intervals describe the retained artifacts separately. This nested design captures search variability without treating task repetitions as repeated evolution.

Analysis and reproducibility. Arm order is randomized within matched tasks, repeated executions use independently recorded randomization, and inference never pools within-task repetitions as independent tasks. Paired differences and task-level intervals follow Section 5; additional confirmatory contrasts share a predefined error budget. The number of independent tasks follows the desired precision and available budget, and is fixed before outcome analysis. Every reported result must identify the dataset partition, all package digests, model configuration, scorer versions, complete search cost, excluded cases, and run artifacts. Search curves use development scores; final claims use the frozen unseen-task comparison.

7 DISCUSSION AND CONCLUSION

7.1 WHAT THE ORGANIZATION RETAINS

An EEO retains procedures and resource relationships that can be applied beyond the execution that motivated them. This differs from keeping a conversation transcript available to the next task: the revision changes the instructions, examples, or domain references that define how the organization works. The transcript supplies evidence for a change; the package is the object that carries it forward. Its exact identity makes that distinction observable. A successful one-off repair with no retained package change does not instantiate this update mechanism.

The package boundary also exposes a substantive design choice. Procedures, role responsibilities, capability references, and supporting files are revised as one attributable object. This makes it possible to ask whether an improvement depends on a particular professional configuration, rather than only whether a new prompt scores better. Packaging alone does not produce that improvement. The matched GEPA control tests whether the particular revision policy adds value on the same substrate. Comparing the fixed generic team with the fixed specialized team isolates specialization; comparing the fixed specialized team with EEO-text isolates retained updating. A direct comparison between the generic team and EEO-text combines both changes.

7.2 WHEN SPECIALIZATION HELPS OR HINDERS

Specialization is most plausibly useful when a task family shares recurring operations and acceptance criteria: source assessment, population reconciliation, computation, revision-specific verification, and evidence-preserving rendering. The package can place these obligations at the roles and handoffs where they matter. The same structure can be costly when the task requires only faithful rewriting of supplied material. Research Studio’s alternative workflow declarations make this variation explicit, but the orchestrator still has to choose an appropriate entry point. Organizational sophistication is therefore not measured by the number of experts invoked.

A retained procedure can improve one failure mode while degrading another. Additional verification may reduce unsupported quantitative claims while increasing delay or abstention. A domain-specific reference can improve familiar tasks and introduce irrelevant assumptions into a new family. Retention evaluation must consequently cover previously useful task types, with the same outcome definitions and recorded resource costs. A mean gain across tasks does not establish that every task or every professional capability improved.

7.3 LIMITS OF THE UPDATE BOUNDARY

Frozen executables and inherited capability references make candidate differences inspectable, but can exclude the tool or algorithm change needed to resolve a failure. Such a failure calls for a separately authorized change to the professional substrate, followed by a new comparison context. It cannot be interpreted as evidence that further text search must eventually solve the problem. Within the allowed space, a text edit can still substantially alter behavior or misuse an existing capability; structural validity does not establish semantic safety.

The evaluation boundary imposes another practical cost. Once outcomes influence candidate selection, they become development evidence. Reusing the same test tasks across successive organizational revisions can make repeated apparent improvement an artifact of adaptation to the test set. Fresh partitions and independent search episodes are needed for the corresponding generalization and method-stability claims. The resource cost of collecting and evaluating that evidence belongs alongside the computational cost of search.

7.4 CONCLUSION

EEO makes reusable professional expertise an explicit organization of roles, procedures, capabilities, and handoffs. Domain specialization establishes its working resources and acceptance contracts; experience-driven revision changes how subsequent tasks are performed. OpenCorvus supplies the package, execution-binding, candidate-validation, and comparative-evidence interfaces that connect these changes. The analysis identifies the full-cost and independent-evaluation conditions under which a retained revision can be assessed as useful. The resulting experimental design separates expert collaboration, professional specialization, and retained updating, and evaluates their effects on future-task quality, resource use, and prior capabilities.

AI USE STATEMENT

This manuscript was drafted and revised with a language-model assistant in Codex. A separate model instance provided critical review. Assistance included source inspection, literature retrieval, mathematical checks, and PDF preparation. The three original method figures were generated with the built-in image-generation tool, documented as GPT Image 2, and inspected for text and scientific consistency. The tool does not expose an exact backend snapshot; the prompts and selected assets accompany the manuscript. These activities are not empirical validation. Human author identities and submission responsibility remain to be confirmed.

REFERENCES

- Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziemis, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning. arXiv:2507.19457, version 2, 2025. URL <https://arxiv.org/abs/2507.19457v2>. Version 2 revised in 2026; cited version is the arXiv manuscript.
- Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963. doi: 10.1080/01621459.1963.10500830. URL <https://doi.org/10.1080/01621459.1963.10500830>.
- Sirui Hong, Mingchen Zhuge, Jiaqi Chen, Xiwu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu, and Jürgen Schmidhuber. MetaGPT: Meta programming for a multi-agent collaborative framework. arXiv:2308.00352, version 7, 2023. URL <https://arxiv.org/abs/2308.00352v7>.
- Shengran Hu, Cong Lu, and Jeff Clune. Automated design of agentic systems. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/file/36b7acf6f6010652b3f2a433774a66fe-Paper-Conference.pdf.
- OpenCorvus contributors. OpenCorvus: Source snapshot and retained project evidence. Software and project records, revision bd2d6bd18eda, 2026. URL <https://github.com/yangheng95/opencorvus/tree/bd2d6bd18eda05f35326adc231289e56d0a2dfcf>. Source inspection; not an independently reproduced experiment.
- Xingyao Wang, Boxuan Li, Yufan Song, Frank F. Xu, Xiangru Tang, Mingchen Zhuge, Jiayi Pan, Yueqi Song, Bowen Li, Jaskirat Singh, Hoang H. Tran, Fuqiang Li, Ren Ma, Mingzhang Zheng, Bill Qian, Yanjun Shao, Niklas Muennighoff, Yizhe Zhang, Binyuan Hui, Junyang Lin, Robert

- Brennan, Hao Peng, Heng Ji, and Graham Neubig. OpenHands: An open platform for AI software developers as generalist agents. arXiv:2407.16741, version 3, 2024. URL <https://arxiv.org/abs/2407.16741v3>.
- Xingyao Wang, Simon Rosenberg, Juan Michelini, Calvin Smith, Hoang Tran, Engel Nyst, Rohit Malhotra, Xuhui Zhou, Valerie Chen, Robert Brennan, and Graham Neubig. The OpenHands software agent SDK: A composable and extensible foundation for production agents. arXiv:2511.03690, version 2, 2025. URL <https://arxiv.org/abs/2511.03690v2>. Version 2 revised in 2026; cited version is the arXiv manuscript.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryen W. White, Doug Burger, and Chi Wang. AutoGen: Enabling next-gen LLM applications via multi-agent conversation. arXiv:2308.08155, version 2, 2023. URL <https://arxiv.org/abs/2308.08155v2>.
- John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. arXiv:2405.15793, version 3, 2024. URL <https://arxiv.org/abs/2405.15793v3>.
- Siyu Yuan, Kaitao Song, Jiangjie Chen, Xu Tan, Dongsheng Li, and Deqing Yang. EvoAgent: Towards automatic multi-agent generation via evolutionary algorithms. In *Proceedings of NAACL: Human Language Technologies (Volume 1: Long Papers)*, pp. 6192–6217. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.naacl-long.315. URL <https://aclanthology.org/2025.naacl-long.315/>.
- Jiayi Zhang, Jinyu Xiang, Zhaoyang Yu, Fengwei Teng, Xiong-Hui Chen, Jiaqi Chen, Mingchen Zhuge, Xin Cheng, Sirui Hong, Jinlin Wang, Bingnan Zheng, Bang Liu, Yuyu Luo, and Chenglin Wu. AFlow: Automating agentic workflow generation. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/file/5492ecbce4439401798dcd2c90be94cd-Paper-Conference.pdf.
- Wanjia Zhao, Mert Yuksekgonul, Shirley Wu, and James Zou. Sirius: Self-improving multi-agent systems via bootstrapped reasoning. arXiv:2502.04780, version 1, 2025. URL <https://arxiv.org/abs/2502.04780v1>.

A IMPLEMENTATION MAPPING AND REPRODUCTION

Source boundary. The inspected implementation is OpenCorvus revision bd2d6bd18eda (OpenCorvus contributors, 2026). The accompanying source map records the full revision identifier, exact files, and symbols. Source inspection establishes declared representations and checked code paths; it does not supply task-success measurements. Historical project claims are not substituted for controlled executions.

Method component	Implementation anchor	Scope
Organization	research-studio/expert-squad.jsonc	Five expert roles, declared capabilities, procedural resources, and report workflow.
Task binding	task-package-revision-binding.ts	Exact package revision recorded for a task.
Candidate space	expert-squad-evolution-integrity.ts	Declared path, graph, inherited-reference, and frozen-file checks.
Candidate publication	publish-evolution-artifact.ts	Identities and differences derived from validated parent/candidate resource sets.
Trial specification	expert-squad-evolution-artifact.ts	Cases, repetitions, scorer and environment identities, partitions, and budgets.
Comparison	Evolution Lab comparison.ts	Matched evidence, score differences, uncertainty summaries, and recommendation.
Retained history	expert-squad-evolution-history.ts	Exact revision and comparison-context identity checks.

Table 3: Source correspondence for the EEO method. Full repository-relative paths are in the distributed source map.

Generator and validator. The candidate-author prompt and Evolution Lab README specify text-only mutation with fixed role and workflow structure. The production integrity comparator accepts a broader manifest-edit space and checks its structural consistency. Equation 7 denotes the proposal distribution under the narrower prompt policy; the EEO-text study additionally checks realized diffs for arm eligibility. Equations 5–6 describe the implemented acceptance relation. The function’s returned mutable-path list is the parent’s text closure, whereas the comparison itself uses the union of parent and candidate closures. This distinction matters when a candidate adds a declared text resource.

Operational comparison statistics. The implementation computes per-scorer paired means, medians, sample variances, and approximate Student- t intervals. It uses tabulated quantiles up to 30 degrees of freedom and a normal approximation beyond that. With fewer than two paired values, a serialized point interval is accompanied by an internal insufficiency condition and cannot support a confidence claim. The aggregate interval combines scaled scorer half-widths in quadrature. Correlated scorers evaluated on the same runs need not satisfy the independence approximation underlying that aggregation. The operational recommendation therefore does not provide the confirmatory guarantee of Equation 11. The experimental protocol uses a predefined primary task-level outcome and separately controlled contrasts.

Retention semantics. The comparison returns inconclusive when required comparison evidence or required visual review is unavailable, or when reviewed findings identify reward hacking. Its selection rule requires a positive aggregate lower endpoint, no identified scorer regression, and no increase in candidate failures. These are operational criteria; retention-set performance is an additional research outcome. A feedback-authored package revision follows an explicit user-acceptance path without an empirical comparison. Version activation is a deployment operation and cannot substitute for evidence that the organization learned.

Execution evidence. Persistent ownership and revision-specific artifact records support the source attribution required by Equation 4. Publication provenance can combine eligible preceding action or

turn evidence with evidence from the current tool invocation. A declared-source read record does not by itself establish that every semantically necessary source was declared, nor that the artifact's claims are correct. Review must identify the exact output revision being assessed. These boundaries motivate independent domain checkers without turning runtime bookkeeping into the main research contribution.

Minimum result bundle. A reproducible campaign contains the parent and every candidate package; development and test partition manifests; all trial slots including failed and rejected attempts; model, workspace, tool, permission, and scorer identities; per-task scores and costs; blinded review instructions and verdicts; and the final selection record. Credentials and private source content are excluded from the public bundle, with access restrictions and reproducible substitutes documented where necessary. Exact sample counts, budget values, and execution commands belong to the frozen campaign record.

Figure and manuscript reproduction. The three figures are original color raster assets generated from the accompanying prompt set and included by small LaTeX wrappers. Their structures, trial cells, and worked example represent declared responsibilities and information flow, not measured performance. The figure record identifies the generation mode, visual references, selected-file hashes, and review decisions. The underlying images contain rasterized labels, so wording changes require regenerating the relevant asset and visually checking it again. The manuscript uses the official ICLR 2027 style files without modifying the downloaded originals. Build instructions, reference verification, and final PDF checks accompany the source.